

# Git for Version Control

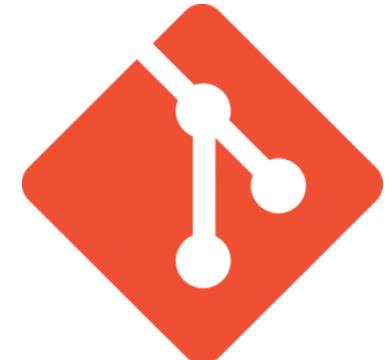
These slides are heavily based on slides created  
by Ruth Anderson.

images taken from <http://git-scm.com/book/en/>

The second part is based on Mike Zastre's  
slides.

# About Git

- Created by Linus Torvalds, creator of Linux, in 2005
  - Came out of Linux development community
  - Designed to do version control on Linux kernel
- Goals of Git:
  - Speed
  - Support for non-linear development (thousands of parallel branches)
  - Fully distributed
  - Able to handle large projects efficiently
  - (*A "git" is a cranky old man. Linus meant himself.*)

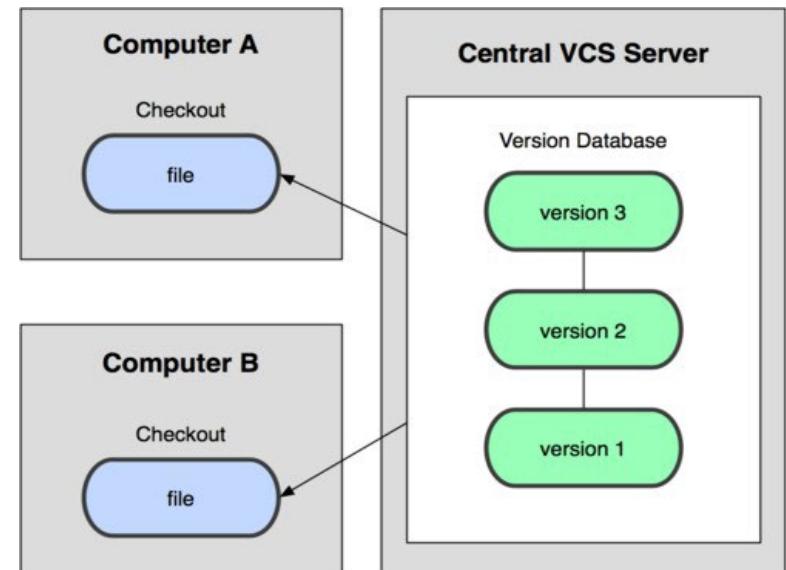


# Installing/learning Git

- Git website: <http://git-scm.com/>
  - Free on-line book: <http://git-scm.com/book>
  - Reference page for Git: <http://gitref.org/index.html>
  - Git tutorial: <http://schacon.github.com/git/gittutorial.html>
  - Git for Computer Scientists:
    - <http://eagain.net/articles/git-for-computer-scientists/>
- At command line: (*where verb = config, add, commit, etc.*)
  - git help verb

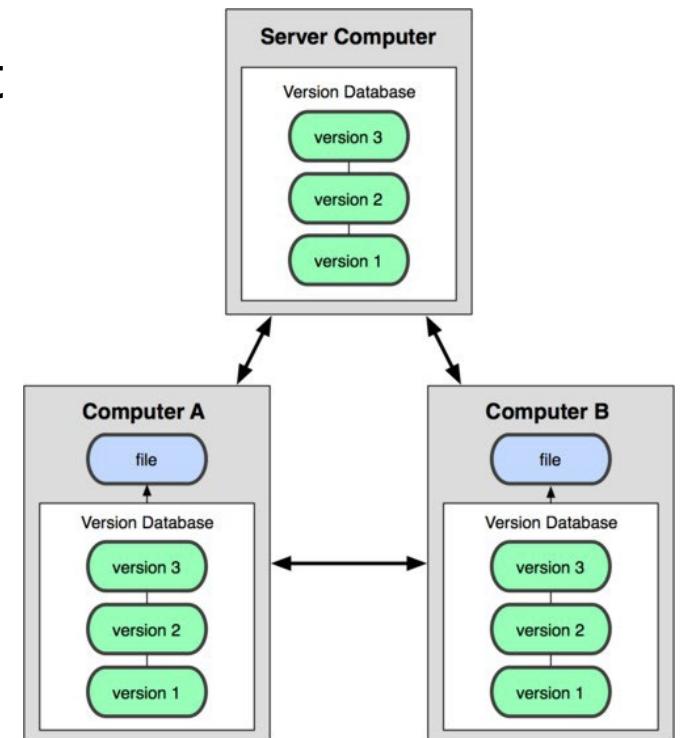
# Centralized VCS

- In Subversion, CVS, Perforce, etc.  
A central server repository (repo)  
holds the "official copy" of the code
  - the server maintains the sole  
version history of the repo
- You make "checkouts" of it  
to your local copy
  - you make local modifications
  - your changes are not versioned
- When you're done, you  
"check in" back to the server
  - your checkin increments the repo's version



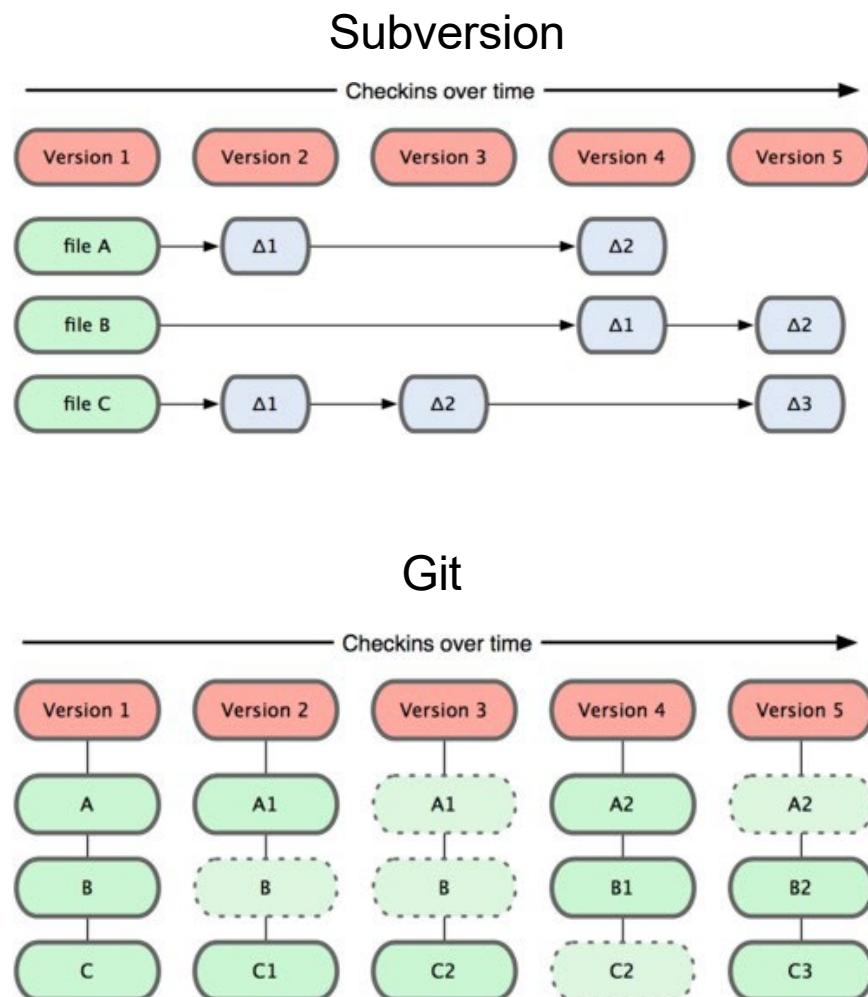
# Distributed VCS (Git)

- In git, mercurial, etc., you don't "checkout" from a central repo
  - you "clone" it and "pull" changes from it
- Your local repo is a complete copy of everything on the remote server
  - yours is "just as good" as theirs
- Many operations are local:
  - check in/out from *local*/repo
  - commit changes to *local*/repo
  - local repo keeps version history
- When you're ready, you can "push" changes back to server



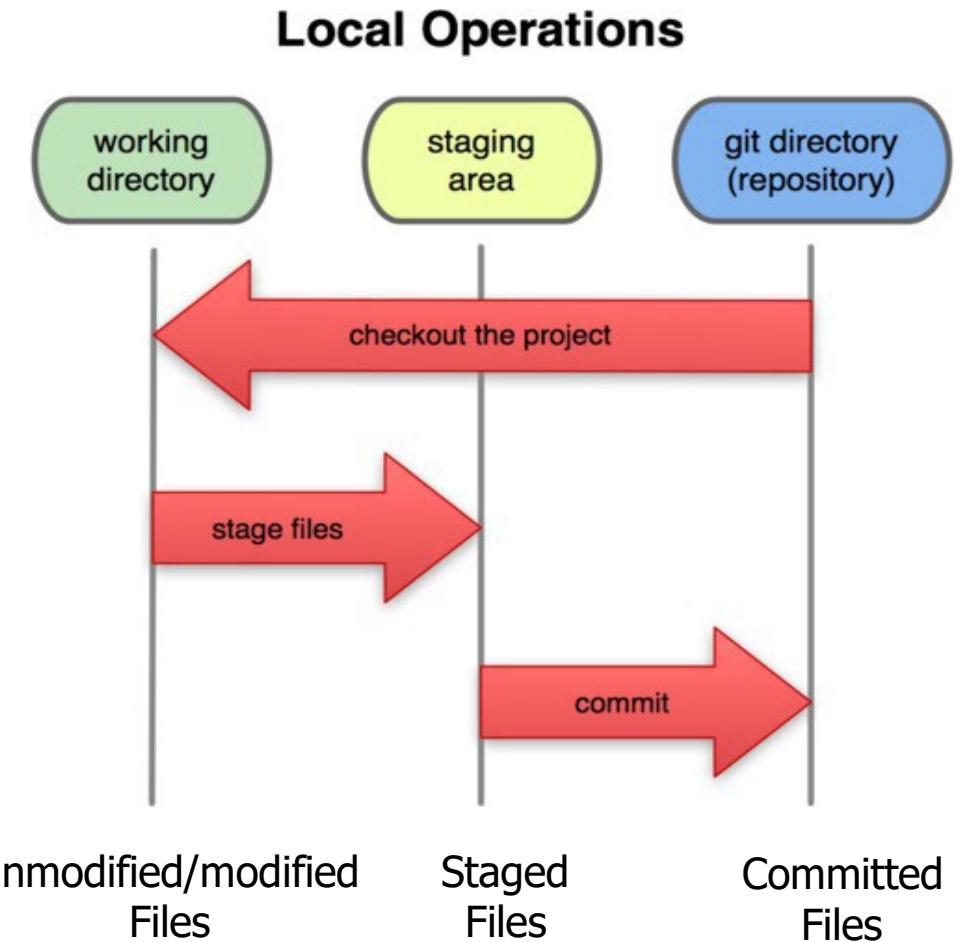
# Git snapshots

- Centralized VCS like Subversion track version data on each individual file.
- Git keeps "snapshots" of the entire state of the project.
  - Each checkin version of the overall code has a copy of each file in it.
  - Some files change on a given checkin, some do not.
  - More redundancy, but faster.



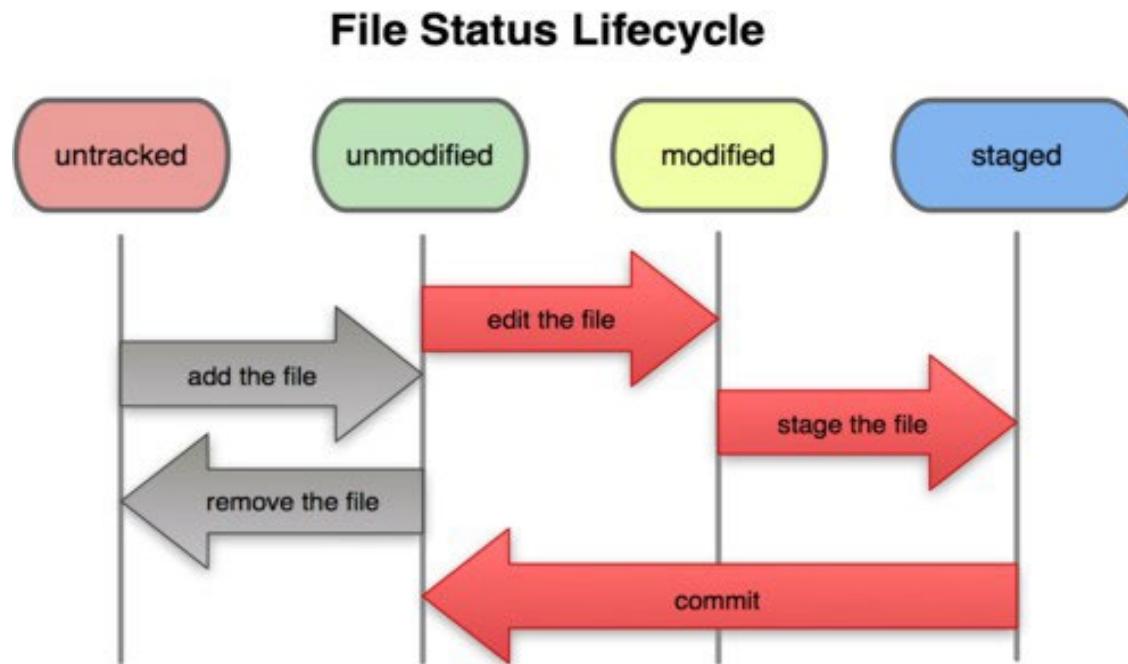
# Local git areas

- In your local copy on git, files can be:
  - In your local repo
    - (committed)
  - Checked out and modified, but not yet committed
    - (working copy)
  - Or, in-between, in a **"staging" area**
    - Staged files are ready to be committed.
    - A commit saves a snapshot of all staged state.



# Basic Git workflow

- **Modify** files in your working directory.
- **Stage** files, adding snapshots of them to your staging area.
- **Commit**, which takes the files in the staging area and stores that snapshot permanently to your Git directory.



# Git commit checksums

- In Subversion each modification to the central repo increments the version # of the overall repo.
  - In Git, each user has their own copy of the repo, and commits changes to their local copy of the repo before pushing to the central server.
  - So Git generates a unique **SHA-1 hash** (40 character string of hex digits) for every commit.
  - Refers to commits by this ID rather than a version number.
  - Often we only see the first 7 characters:
    - 1677b2d Edited first line of readme
    - 258efa7 Added line to readme
    - 0e52da7 Initial commit

# Initial Git configuration

- Set the name and email for Git to use when you commit:
  - `git config --global user.name "Bugs Bunny"`
  - `git config --global user.email bugs@gmail.com`
  - You can call `git config -list` to verify these are set.
- Set the editor that is used for writing commit messages:
  - `git config --global core.editor nano`
    - (it is vim by default)

# Creating a Git repo

*Two common scenarios: (only do one of these)*

- To create a new **local Git repo** in your current directory:
  - `git init`
    - This will create a `.git` directory in your current directory.
    - Then you can commit files in that directory into the repo.
  - `git add filename`
  - `git commit -m "commit message"`
- To **clone a remote repo** to your current directory:
  - `git clone url localDirectoryName`
    - This will create the given local directory, containing a working copy of the files from the repo, and a `.git` directory (used to hold the staging area and your actual local repo)

# Git commands

command	description
<code>git clone <i>url</i> [<i>dir</i>]</code>	copy a Git repository so you can add to it
<code>git add <i>file</i></code>	adds file contents to the staging area
<code>git commit</code>	records a snapshot of the staging area
<code>git status</code>	view the status of your files in the working directory and staging area
<code>git diff</code>	shows diff of what is staged and what is modified but unstaged
<code>git help [<i>command</i>]</code>	get help info about a particular command
<code>git pull</code>	fetch from a remote repo and try to merge into the current branch
<code>git push</code>	push your new branches and data to a remote repository
others: <code>init</code> , <code>reset</code> , <code>branch</code> , <code>checkout</code> , <code>merge</code> , <code>log</code> , <code>tag</code>	

# Add and commit a file

- The first time we ask a file to be tracked, *and every time before we commit a file*, we must add it to the staging area:
  - `git add Hello.java Goodbye.java`
    - Takes a snapshot of these files, adds them to the staging area.
    - In older VCS, "add" means "start tracking this file." In Git, "add" means "add to staging area" so it will be part of the next commit.
- To move staged changes into the repo, we commit:
  - `git commit -m "Fixing bug #22"`
  - All these commands are acting on your local version of repo.

# Viewing/undoing changes

- To view status of files in working directory and staging area:
  - git status                          or git status -s (short version)
- To see what is modified but unstaged:
  - git diff
- To see a list of staged changes:
  - git diff --cached
- To see a log of all changes in your local repo:
  - git log                          or git log --oneline (shorter version)  
1677b2d Edited first line of readme  
258efa7 Added line to readme  
0e52da7 Initial commit
    - git log -5 (to show only the 5 most recent updates), etc.

# Interaction w/ remote repo

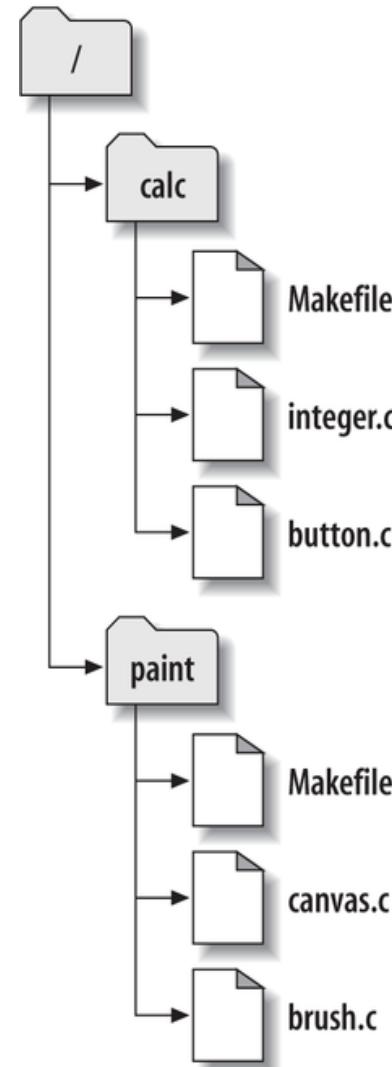
- **Push** your local changes to the remote repo.
- **Pull** from remote repo to get most recent changes.
  - (fix conflicts if necessary, add/commit them to your local repo)
- To fetch the most recent updates from the remote repo into your local repo, and put them into your working directory:
  - `git pull origin master`
- To put your changes from your local repo in the remote repo:
  - `git push origin master`

# GitHub and GitLab

- [GitHub.com](https://GitHub.com) is a site for online storage of Git repositories.
  - You can create a **remote repo** there and push code to it.
  - Many open source projects use it, such as the Linux kernel.
  - You can get free space for open source projects, or you can pay for private projects.
    - Free private repos for educational use
- [GitLab.com](https://GitLab.com) is also a site for online storage of Git repositories.
  - But it also allows educational institutions to create their GitLab sites.
  - Example: [gitlab.csc.uvic.ca](https://gitlab.csc.uvic.ca)
- *Question:* Do I always have to use GitHub or GitLab to use Git?
  - *Answer:* No! You can use Git locally for your own purposes.
  - Or you or someone else could set up a server to share files.
  - Or you could share a repo with users on the same file system, as long everyone has the needed file permissions).

# Basic Concepts: Working Copy

- A **working copy / local repository** is an ordinary directory on your local system
- These are the files you edit
- When changes are at some suitable stage, add & commit your changes in your local repo
  - In practice, this usually means metadata files in the .git directory within the working copy are added or changed (.git directory not shown)
- You may even decide to push committed changes to the remote repo



# Basic Concepts: Working Copy

- Obtaining a working copy means either **cloning** an existing repository or performing **init** in some existing directory
  - This is normally done **only once** per working copy
  - All work is done in a (you guessed it!) working copy
  - Working copy is simply a set of directories & files
- Repository access methods differ:
  - direct access via local disk (file:/// **(ugh!)**)
  - via ssh:// (when security is very important)
  - http:// or https:// (as used by GitHub and GitLab)
  - original access method for working copy / local repo is stored as metadata in .git subdirectories in that working copy)

Please do not use **git init** for this course!

# Basic Concept: Working Copy

- Example: get working copy of the "calc" project from a (made up!) git.example.com

```
$ git clone ssh://stewie@git.example.com/repos/calc
Cloning into 'calc'...
<password for stewie>
remote: Counting objects: 37, done.
remote: Compressing objects: 100% (31/31), done.
remote: Total 37 (delta 5), reused 0 (delta 0)
Receiving objects: 100% (37/37), done.
Resolving deltas: 100% (5/5), done.
Checking connectivity... done.
```

```
$ cd calc
$ git ls-files # could even use `tree` here...
Makefile
button.c
integer.c

$
```

# Basic Concepts: Commit

- Suppose you wish to keep track of some changes to **button.c**
  - You edit the file using your normal workflow
  - Time and date on edited file will be more recent than time and date of file in local repo
  - Changes are recorded by committing your changed file to the repository
- We first stage our changes (via **add**) ...
- ... and then make a "permanent" record of the change (**commit**)

```
$ pwd  
calc
```

```
$ git add button.c  
$ git commit -m "Fixed the geometry of button for v3 of library"  
[master 12788ce] Fixed the geometry of button for v3 of library  
 1 file changed, 1 deletion(-)  
$
```

# Basic Concepts: Commit

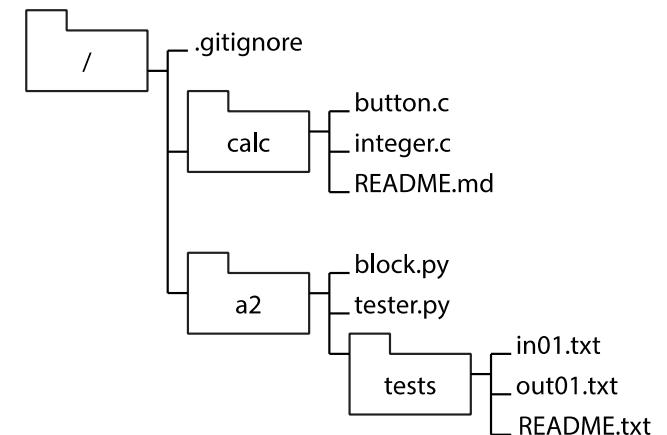
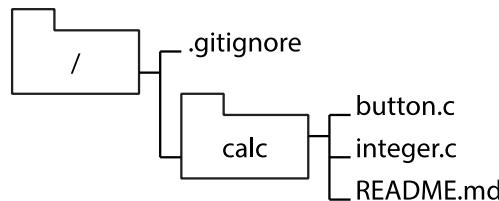
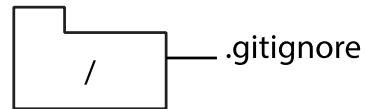
- Note that commits are **always** to our local repo
  - i.e., take place within our working/local copy
- We can have a sequence of **many** such commits as we work through sets of changes
- **To have the results of the commits available to others...**
  - ... or to ourselves on a different machine ...
  - ... we must push them to a remote repository
- Observations:
  - Commits can be quite frequent (if needed by our workflow)
  - **Pushes are much less frequent** (with such a push perhaps reflecting our work has reached some suitable state and is ready for other team members to pull).
  - Example: writing and completing each assignment will require several commits
  - Example: submitting an assignment will require at least one push

# Basic Concepts: Commit

- Each **commit** results **in the creation of a new snapshot of the contents** in our working copy / local repository
- Snapshots are kept in chronological order
  - **git log** produces a list of the snapshots
  - Default log output order is reverse chronological (i.e., most recent commit/snapshot is listed first)
- **git status** reports the relationship amongst files in our working directory with what within git's local repository
  - More precisely, "status" tells us what has changed in our working directory...
  - ... and therefore what may need to be "add"ed and "commit"ed to the local repository

# Basic Concepts: Commit

time



HEAD

18431a0

HEAD

18431a0

3da5a35

HEAD

3da5a35

85687aa

initial  
commit

commit after adding "calc"  
with its files

commit after adding "a2"  
with its files and  
subdirectory

# (steps leading to third commit)

```
$ pwd  
/home/stewie/project
```

```
$ git add * # Let git know what files are to be committed; recursive on directories
```

```
$ git status  
On branch master  
Your branch is ahead of 'origin/master' by 1 commit.  
(use "git push" to publish your local commits)
```

```
Changes to be committed:  
(use "git reset HEAD <file>..." to unstage)
```

```
new file: a2/block.py  
new file: a2/tester.py  
new file: a2/tests/README.txt  
new file: a2/tests/in01.txt  
new file: a2/tests/out01.txt
```

```
$ git commit -m "Now the A#2 directory (a2) is in place" # Note message  
[master 85687aa] Now the A#2 directory (a2) is in place  
5 files changed, 6 insertions(+)  
create mode 100644 a2/block.py  
create mode 100644 a2/tester.py  
create mode 100644 a2/tests/README.txt  
create mode 100644 a2/tests/in01.txt  
create mode 100644 a2/tests/out01.txt
```

# Basic Concepts: Commit

```
$ git log
```

```
commit 85687aa056e299897153a3125c1826f64581bdc5
```

```
Author: Stewie Griffin <stewie@uvic.ca>
```

```
Date: Thu Apr 15 10:05:54 2018 -0700
```

Now the A#2 directory (a2) is in place

```
commit 3da5a353956c320fbe8e585cd692b173e44b06c1
```

```
Author: Stewie Griffin <stewie@uvic.ca>
```

```
Date: Thu Apr 15 10:01:57 2018 -0700
```

Added calc and some files

```
commit 18431a0b85f0645c98e4cce311074594a19a38d
```

```
Author: Stewie Griffin <stewie@uvic.ca>
```

```
Date: Thu Apr 15 09:38:34 2018 -0700
```

Initial commit (just .gitignore for now)

```
$
```



# Basic Concepts: Commit

```
$ git status
```

On branch master

Your branch is ahead of 'origin/master' by 2 commits.  
(use "git push" to publish your local commits)

nothing to commit, working directory clean

```
$ git push
```

<verbiage>

<password for stewie>

Counting objects: 13, done.

Delta compression using up to 8 threads.

Compressing objects: 100% (10/10), done.

Writing objects: 100% (12/12), 1.11 KiB | 0 bytes/s, done.

Total 12 (delta 0), reused 0 (delta 0)

To ssh://stewie@git.example.com/project

  18431a0..85687aa master -> master

```
$ git status
```

On branch master

Your branch is up-to-date with 'origin/master'.

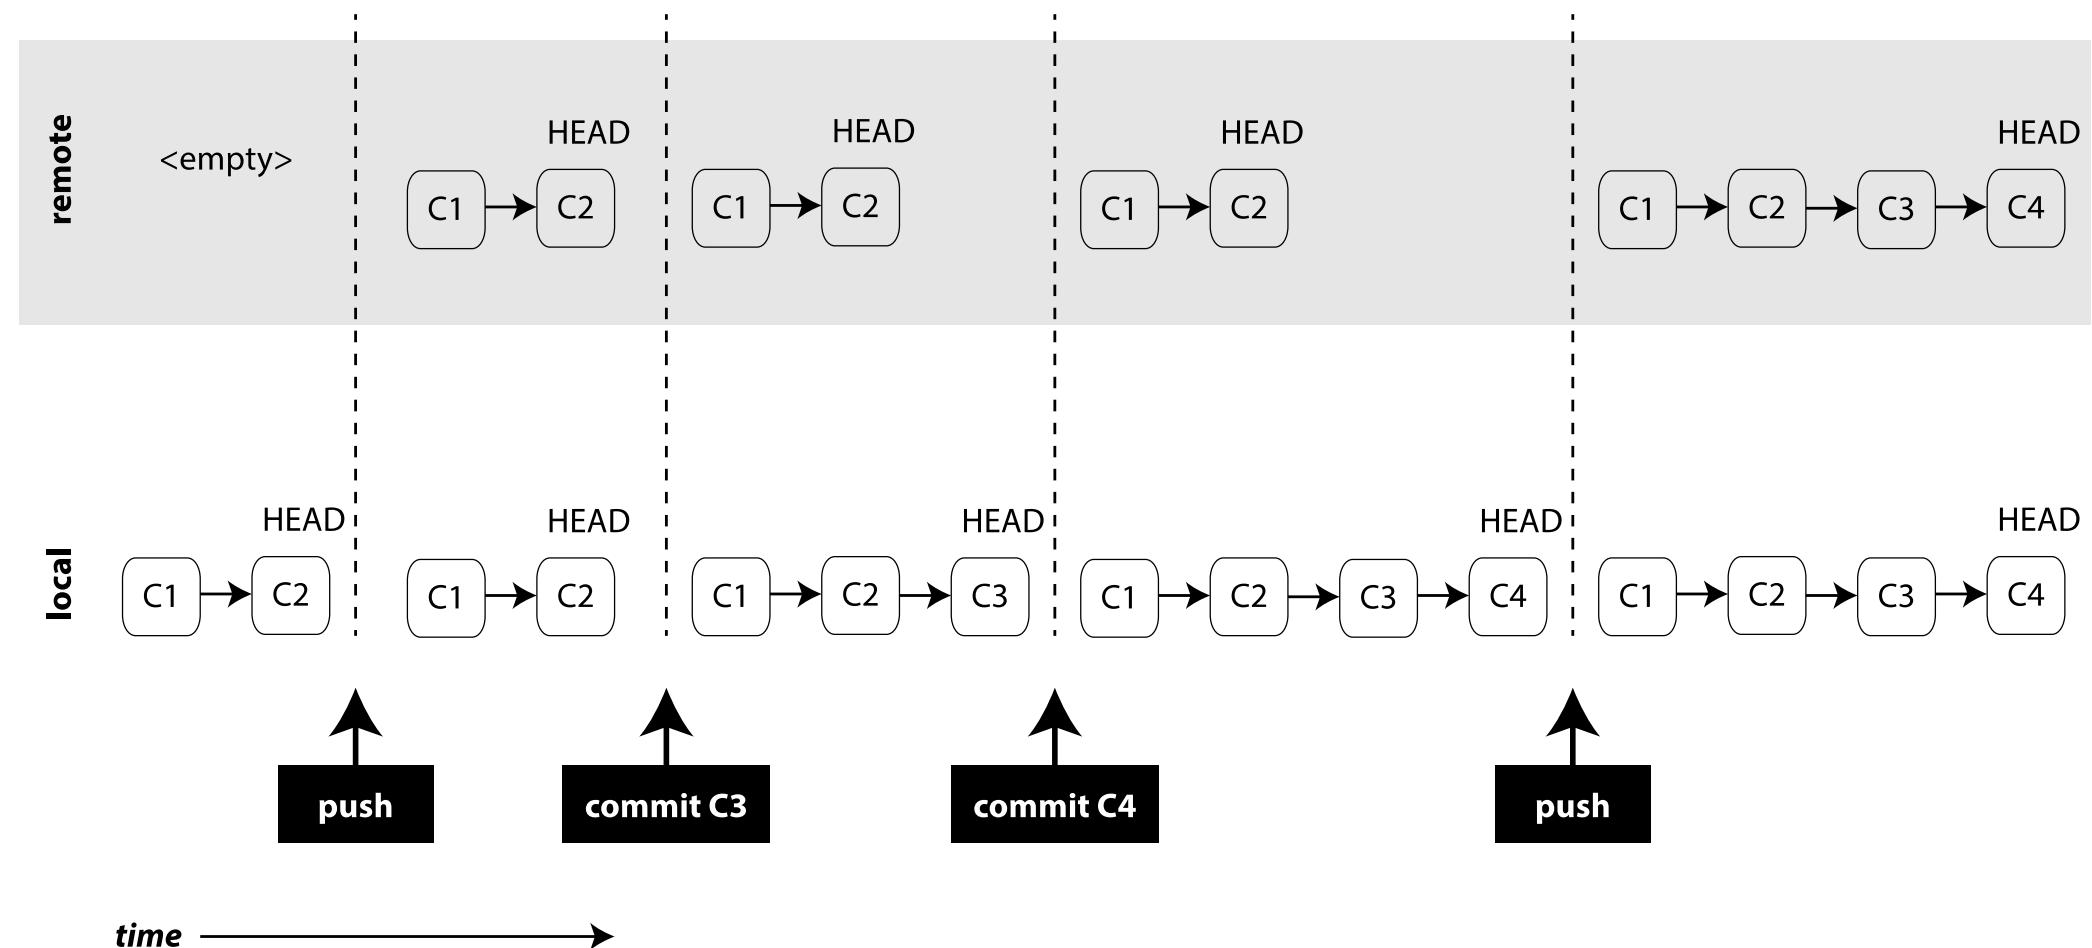
nothing to commit, working directory clean

```
$
```

# commit != push

- git is different from many other kinds of VCSes
  - commits are made to the local repo, **not the remote**
  - making these changes available to others means transferring data from the local repo to the remote repo (i.e., a push)
- git separates **the tracking of file/directory changes** from **their storage on remote servers/repositories**
  - this is very different from Subversion, Perforce, etc.
  - ... and can seem a bit confusing on first encounter

# commit != push



# Basic Concepts: Update

- What if Meg starts working on the project after someone else's commit?
  - Assume she made a working copy of Stewie's repo some time ago
  - (Also assume here she has read & write privileges on Stewie's remote.)
  - Let's also assume she was not working on button.c in calc
- She can ask git to bring her working copy "up to date"
  - git will only update files for which there are changes on the remote
  - Principle: Make sure to update often if working with a group on a project that uses a repository!

```
$ pwd  
/home/meg/calc
```

```
$ git pull  
<password for meg>  
remote: Counting objects: 7, done.  
remote: Compressing objects: 100% (4/4), done.  
<... snip ...>  
Fast-forward  
 calc/button.c | 1 +  
 1 file changed, 1 insertion(+)
```

```
$
```

# Basic Concepts: Update

- "git pull" is actually two commands together:
  - "git fetch" followed by "git merge"

```
$ pwd  
/home/meg/calc
```

```
$ git fetch  
<password for meg>  
remote: Counting objects: 7, done.  
<snip>
```

```
$ git log --name-status  
commit eb6c8a6ffb2e2a70a89e4a89db8d62b5a22eccd4  
Author: Stewie Griffin <stewie@uvic.ca>  
Date: Thu Apr 15 10:27:23 2018 -0700  
    Added headers so buttons can be beveled
```

```
M      calc/button.c
```

# Basic Concepts: Update

```
$ pwd  
/home/meg/calc
```

```
$ git merge origin/master  
Updating 85687aa..eb6c8a6  
Fast-forward  
 calc/button.c | 3 +++  
 1 file changed, 3 insertions(+)  
$
```

- Therefore "git pull" is the same as the following two commands in succession
  - git fetch
  - git merge origin/master
- We used "git log --name-status" to obtain the names of files that are different from our working copy and the remote repo