

Regular Expressions

- Background
- Sets of strings
- Stating a regular expression (simple)
- Python `re` module (simple)
- A bit of theory

String patterns

- We all use searches where we provide strings or substrings to some module or mechanism
 - Google search terms
 - Filename completion
 - Command-line wildcards
 - Browser URL completion
 - Python string routines `find()`, `index()`, etc.
- Quite often these searches are simply expressed as a particular pattern
 - An individual word
 - Several words where some are strictly required while some are not
 - The start or end of particular words -- or perhaps just the string appearing within a larger string
- This works well if strings follow the format we expect...

String patterns

- Sometimes, however, we want to express a more complex pattern
 - The set of all files ending with either `.c` or `.h`
 - The set of all files starting with `in` or `out`.
 - The set of all strings in which `FREQ` appears as a string (but not `FREQUENCY` or `INFREQUENT`, but `fReQ` is fine)
 - The set of all strings containing dates in `MM/DD/YYYY` format.
- Such a variety of patterns used to require language-specific operations
 - SNOBOL
 - Pascal
- More troubling was that most non-trivial patterns required several lines of code to express (i.e., a series of "if-then-else" statements)
 - This is a problem as the resulting code can obscure the patterns for which we are searching
 - Even worse, changing the pattern is tedious and error-prone as it means changing the structure of already written code.

C code to check for DD/MM/YYYY format

```
int is_date_format(char *check) {  
  
    if (!isdigit(check[0]) || !isdigit(check[1])) {  
        return 0;  
    }  
  
    if (!isdigit(check[3]) || !isdigit(check[4])) {  
        return 0;  
    }  
  
    for (i = 6; i < 10; i++) {  
        if (!isdigit(check[i])) {  
            return 0;  
        }  
    }  
  
    if (check[2] != '/' || check[5] != '/') {  
        return 0;  
    }  
  
    /* Still haven't even figured out if the DD makes sense, let alone  
     * the MM!!!!  
     */  
  
    return 1;  
}
```

Regular expressions

- Needed: a language-independent approach to expressing such patterns
- Solution: a **regular expression**
 - Sometimes called a **regex** or **regexp**
- They are written in a formal language and have the property that we can build very fast recognizers for them
- Part of a hierarchy of languages
 - Type 0: unrestricted grammars
 - Type 1: context-sensitive grammars
 - **Type 2: context-free grammars**
 - **Type 3: regular grammars**
- Type 2 and 3 grammars are used in Computer Science
 - Type 2 is used in parsers for computer languages (i.e., compilers)
 - Type 3 is used in regular expressions and lexical analyzers for compilers

grep

- We already can use regular expressions in Unix at the command line
- The grep utility accepts two sets of arguments
 - **grep: global regular expression print**
 - argument 1: A regular expression
 - argument 2: A set of files through which grep will try to find strings matching the regex
- The syntax for a regex is grep is somewhat similar to what we will use in Python
 - grep is a Very Old Tool (i.e., from 1973)
 - superseded somewhat in Modern Times by fgrep (fixed-string grep)
 - a variety of extensions, optimizations, etc. exist
- Example: search for variants on apple

grep

```
apple  
apples  
Apple Pie  
APPLE SUX!  
apple-  
apple-fruit  
"Apple is the greatest!"  
My best friend is an apple.  
pineapple  
Crabapple  
fruit-apple
```

contents of fruitstuff.txt

```
unix$ grep -i ^apple fruitstuff.txt  
apple  
apples  
Apple Pie  
APPLE SUX!  
apple-  
apple-fruit
```

```
unix$ grep apple fruitstuff.txt  
apple  
apples  
apple-  
apple-fruit  
My best friend is an apple.  
pineapple  
Crabapple  
fruit-apple
```

```
unix$ grep ^a.ble fruitstuff.txt  
apple  
apples  
apple-  
apple-fruit
```

```
unix$ grep -w apple fruitstuff.txt  
apple  
apple-  
apple-fruit  
My best friend is an apple.  
fruit-apple
```

```
unix$ grep apple$ fruitstuff.txt  
apple  
pineapple  
Crabapple  
fruit-apple
```

More general regular expressions

- Our grep examples were relatively simple
- Sometimes we want to denote more complex sets of strings
 - strings where the beginning and end match a pattern, while everything in-between can vary
 - all possible spellings of a particular name
 - match non-printable characters
 - catch possible misspellings of a particular word
 - match Unicode code points
- And we may want even more:
 - when matching patterns to strings, extract the actual match itself
 - look for strings where the matched pattern repeats exactly later in the same string
 - extract multiple matches from one string

Metasymbols

- Fully-fledged regexes initially look intimidating because of the metasymbols
- However, all that is required to understand them is patience
- Regexes never loop...
- ... nor are they ever recursive
- Understanding them means reading from left-to-right!
- However, first some metasymbols

symbol/example	meaning
.	match any char except \n
a*	zero or more reps of 'a'
a+	one or more reps of 'a'
a?	zero or one rep of 'a'
a{5}	exactly 5 reps of 'a'
a{3,7}	3 to 7 reps of 'a'
[abc]	any one character in the set {'a', 'b', 'c'}
[^abc]	any one character not in the set of {'a', 'b', 'c'}
a b	match 'a' or 'b'
(...)	group a component of symbols in the regex
\	escape any metasymbol (caution!)

Special pattern elements

symbol	meaning
\d	Any decimal digit character
\w	Any alphanumeric character
\s	Any whitespace character (\t \n \r \f \v)
\b	Empty string at a word boundary
^	match 0 characters at the start of the string
\$	match 0 characters at the end of the string
\D	match any non-digit character (opposite of \d)
\W	match any non-alphanumeric character (opposite of \w)
\S	match any non-whitespace character
\B	empty string (i.e., 0 characters) not at a word boundary
\number	matches text of group number

Python regular expressions

- The `re` module
 - Introduced into Python in version 1.5
 - (Don't use the `regex` module which is an older release of a regular-expression library)
 - Used to be slower than regex, but is now as fast if not faster
 - Supports **named groups**
 - Supports **non-greedy matches** (we'll cover this later)
- Note:
 - Regular expression syntax is generally the same from language to language and library to library (e.g., Python, Perl, Ruby)
 - However, sometimes there are differences in the way some features are expressed (e.g., groups, escaped characters)
 - Whenever you move to different implementations, always have the library reference nearby.

Simple example

```
>>> text1 = 'Hello spam...world'
>>> text2 = 'Hello spam...other'

>>> import re
>>> matchobj = re.match('Hello.*world', text2)
>>> print (matchobj)
None

>>> if re.match('Hello.*world', text2):
...     print ("It's the end of the world")
... else:
...     print ("The end of the world is nowhere in sight")
...
The end of the world is nowhere in sight
>>
```

Previous example

- The regular-expression match was applied to string `text2`
 - Regex specified a string with "Hello" followed by 0 or many characters followed by "world"
 - The match did not succeed, therefore the value `None` was returned
 - In Python, `None` may be used as part of a conditional expression (i.e., has similar meaning to "`False`".)
- Even though the name of the RE method was `match()`, we did not use any syntax to extract out some result of the match
 - Which is just as well as there was no match.
 - However, if we wanted to extract out some result, we must use parentheses.
- Let's look at the example again, but this time include the other string in our use of `match`
 - Note that in the following example the "`import re`" is left out (i.e., we assume it was executed earlier in the session)

Simple example

```
>>> text1 = 'Hello spam...world'
>>> text2 = 'Hello spam...other'

>>> matchobj = re.match('Hello(.*)world', text1)
>>> print (matchobj)
<_sre.SRE_Match object at 0x10043b8a0>

>>> hello_list = [text1, text2]
>>> for t in hello_list:
...     matchobj = re.match('Hello(.*)world', t)
...     if matchobj:
...         print (t, " --> match --> ", matchobj.group(1))
...     else:
...         print (t, " --> no matches")
...
Hello spam...world  --> match -->  spam...
Hello spam...other  --> no matches
```

Previous example

- The match did succeed when applied to **text1**
 - The result is a match object
 - This has an interface which is used to extract matched substrings
 - In this case, we extracted the substring matching the pattern in the parentheses
- The parameter passed to **group** corresponds to the order of left parenthesis
 - A regular expression can have several such groups given the use of parentheses
 - Groups can even be nested (i.e., nested parentheses)...
 - ... but **they can never overlap.**
 - Programmers make extensive use of groups in regular expressions
 - It helps make code more robust and less dependent on an exact format.



Speed concerns

- So far we have specified the regular expression for every use of an `re` operation
- For occasional regex matching this is fine
- However, each time the match is performed the Python interpreter must re-interpret the regex
 - This means the regex must be re-parsed and the state machine re-constructed.
 - If we want to search many strings using the same regex, it makes sense to eliminate the overhead of repeating this work.
 - To eliminate the repeated work, we must **compile the pattern**

Speed concerns

- When using this style of regex matching, we work with a pattern object
 - Resulting code is much, much faster
 - Note, however, the compilation itself takes up some cycles.
- For now, just be aware there exist the two styles of invoking re operations
 - One directly specifying the regex in call to `match()`, `search()`, etc.
 - The other using a pattern object returned from `re.compile()` for which we call `match()`, `search()`, etc.).

Regex as a state machine



Compiled pattern

```
>>> pattobj = re.compile('Hello(.*?)world')
>>> matchobj = pattobj.match(text1)
>>> print (matchobj)
<_sre.SRE_Match object at 0x10043b8a0>

>>> hello_list = [text1, text2]
>>> for t in hello_list:
...     matchobj = pattobj.match(t)
...     if matchobj:
...         print (t, " --> match --> ", matchobj.group(1))
...     else:
...         print (t, " --> no matches")
...
Hello spam...world  --> match -->  spam...
Hello spam...other  --> no matches
```

Lots in the `re` module

- Python's `re` module has methods for:
 - matching (i.e., finding a match that must start at the beginning of the string)
 - searching (i.e., finding a match that may occur anywhere in the string)
 - substituting
 - precompiling
 - splitting
 - iterating through matches
- Match objects also have several methods
 - We've already seen `group()`
 - There are also `groups()`, `groupdict()`
- Let us look at a few examples, this time with a few more metasymbols included

More complex pattern

```
>>> datetime1 = "20220211T110000"
>>> datetime2 = "20211225T000000Z"
>>> datetime3 = "11/06/2022"
>>> datetime4 = "22/4/14"

>>> matchobj = re.match("(\\d{4})(\\d{2})(\\d{2})T.*", datetime1)
>>> if matchobj:
...     (year, month, day) = matchobj.groups()
...     print (year, month, day)
... else:
...     print ("Error")
...
2022 02 11
```

More complex pattern

```
>>> dates = ["20220211T110000", "20211225T000000Z", "11/06/2022",
             "22/4/14"]

>>> pattobj = re.compile( "(\d\d?)/(\d\d?)/(\d\d(\d\d)?)" )

>>> for d in dates:
...     matchobj = pattobj.match(d)
...     if matchobj:
...         (day, month, year, _) = matchobj.groups()
...         print "%4d%02d%02d" % (int(year), int(month),
int(day))
...     else:
...         print (d, "doesn't match")
...
20220211T110000 doesn't match
20211225T000000Z doesn't match
20220611
140422
```

Another pattern

```
>>> line1 = ".LM +5"
>>> line2 = ".LM filled"
>>> line3 = ".LM 10x"
>>> line4 = ".LM 22"
>>> lines = [line1, line2, line3, line4]

>>> for line in lines:
...     matchobj = re.match("\.LM (\d+)\s*$", line)
...     if matchobj:
...         values = matchobj.groups()
...         print (line, ": matches with value ", values[0])
...     else:
...         print (line, ": DOESN'T match")

.LM +5 : DOESN'T match
.LM filled : DOESN'T match
.LM 10x : DOESN'T match
.LM 22 : matches with value  22
```

Notes from previous example

- Although grouping may be used to control the regular-expression match, not all results need to be extracted
 - Notice that sometimes part of the extracted matches is ignored
 - Always be aware the extracted matches are indexed by opening left parenthesis (i.e., not by your intent as a programmer to extract out particular parts of the match)
- There is often more than one way to phrase the same regular expression
 - Note that "\d\d" is the same as "\d{2}"
 - Which one is better? Depends perhaps on style of programmer, amount of change expected with code, etc. etc.

Variety

- Sometimes our needs vary when working with regexes
 - Sets of strings may be best expressed by alternative strings
 - Regexes may need to be carefully crafted sets of characters
 - Matches may sometimes be required on word boundaries
 - Sometimes all we want is the starting location of the match.
- Python string rules can sometimes interfere with regular expressions
 - The problem is with backslashes
 - Sometimes you must double-up on them (e.g., "\\")

Using search()

```
>>> pattern, string = "A.C.", "xxABCDxx"
>>> matchobj = re.search(pattern, string)
>>> if matchobj:
...     print (matchobj.start())
...
2
>>> pattobj = re.compile("A.*C.*")
>>> matchobj = pattobj.search("xxABCDxx")
>>> if matchobj:
...     print (matchobj.start())
...
2

>>> print (re.search(" *A.C[DE][D-F][^G-ZE]G\t+ ?", "...ABCDEFG\t..").start())
2
>>> print (re.search("A|XB|YC|ZD", "...AYCD..").start())
2
>>> print (re.search("\bABCD", "..ABCD").start())
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
AttributeError: 'NoneType' object has no attribute 'start'
>>> print (re.search(r"\bABCD", "..ABCD").start())
2
>>> print (re.search(r"ABCD\b", "..ABCD").start())
2
```

Problem Solving

- We have seen a variety of **metasymbols** (sometimes referred to as **metacharacters**)
 - Most of them match one or more characters
 - Some, however, are meant to catch a particular position (i.e., they catch zero characters!)
- The simplest positional symbols are `^` and `$`
 - `^`: match beginning of string
 - `$`: match end of string
 - Note that `re.match("<pattern>", string)` is **exactly** the same as `re.search("^<pattern>", string)` if `string` is not multiline
- Another positional symbol is `\b`
 - Matches a word boundary (i.e., zero characters)
 - That is, it matches the position in between characters (one of which is a word character, the other a non-word character)
 - Word characters: `[a-zA-Z0-9_]`
- **Problem 1: Match the word "Chris" in a string, but not "Christmas", "Christine", etc.**

Problem Solving

```
#!/usr/bin/env python3

import re

lines = ['''I said to Chris, "Hey, watch out!"''',
         '''I'll be home for Christmas!''',
         'Christine Faulkner',
         'Chris Flynn',
         'Evert, Chris']

for li in lines:
    if (re.search(r'\bChris\b', li)):
        print (li)
```

```
$ ./prob01.py
I said to Chris, "Hey, watch out!"
Chris Flynn
Evert, Chris
```

Problem Solving

- Words need not be textual
 - They can also be numerical
 - Key point is that non-word characters are neither numbers nor letters (nor the underscore)
- Sometimes we are interested in the shape of number sequences
 - Course numbers
 - Room numbers
 - Serial numbers, product codes, etc.
- **Problem 2: Extract the last four digits from a North American phone number**
 - May be of the form "250-472-5000"...
 - or "250 472 5000"...
 - or "472-5000"
 - or perhaps "250.472.5000" or "+1 250 472 5000"

Problem Solving

```
#!/usr/bin/env python3

import re

lines = ["250-472-5000", "472-5000", "250.472.5000", \
    "+1 250 472 5000", "011 49 9602 4241", "2504725000", \
    "mom's number", "12345 678 90"]

for l in lines:
    matchobj = re.search(r"(\b\d{3}\b[- \.])?\b\d{3}\b[- \.](\b\d{4}\b)", l)
    if matchobj:
        matchobj = re.search(r"\b(?:\d{3})?\d{3}(\d{4})\b", l)
        if matchobj:
            print(l, "-->", matchobj.group(1))
```

```
$ ./prob02.py
250-472-5000 --> 5000
472-5000 --> 5000
250.472.5000 --> 5000
+1 250 472 5000 --> 5000
2504725000 --> 5000
```

Notice the "?:" used in the second search. It makes a set of parentheses "non-matching" but still useful to structure the regex. That is why the group number is still 1 even though the match we want is denoted by the second left-parenthesis

Problem Solving

- We can also use regexes to verify that the format provided as input matches what we expect
 - Example: Input string is in "DD/MM/YYYY" or "MM/DD/YYYY" format
 - Example: String provided is a URI (i.e., proper sets of characters)
- **Problem 3: Obtain a temperature (assumed to be Celsius) and return the number in Fahrenheit**
 - Number is to be an integer
 - There must be only one number in the string
 - No other characters (such as "C") should be at the end
 - $\text{Fahrenheit} = (\text{Celsius} * 9 / 5) + 32$

Input validation

```
#!/usr/bin/env python3

import re

celsius = input("Enter a temperature in Celsius: ")
celsius = celsius.rstrip("\n")

matchobj = re.search(r"^[0-9]+$", celsius) # same as re.match("\d+",...)
if matchobj:
    celsius = int(celsius)
    fahrenheit = (celsius * 9 / 5) + 32
    print ("%d C is %d F" % (celsius, fahrenheit))
else:
    print ("Expecting a number, so I don't understand", celsius)
```

```
$ ./prob03.py
Enter a temperature in Celsius: 30
30 C is 86 F
```

Problem Solving

- However, we should do a bit more
 - The problem statement is perhaps a bit too restrictive.
 - Negative temperatures cannot be given as values.
 - Decimal temperatures also cannot be provided
- The regular expression should accept these
 - And the other code changed to suit (i.e., use `float()` instead of `int()`)

Input validation

```
#!/usr/bin/env python3

import re

celsius = input("Enter a temperature: ")
celsius.rstrip("\n")

matchobj = re.search(r"^-?[0-9]+(\.[0-9]*)?", celsius)
if matchobj:
    celsius, _ = matchobj.groups()
    celsius      = float(celsius)

    fahrenheit = (celsius * 9 / 5) + 32
    print ("%.2f C is %.2f F" % (celsius, fahrenheit))
else:
    print ("Expecting a number, so I don't understand", celsius)
```

```
./prob03.py
Enter a temperature in Celsius: 12.2
12.20 C is 53.96 F
```

Problem Solving

- Our little script could be even more general
 - Rather than just convert from celsius to fahrenheit, it could convert the other direction
 - The starting value can be indicated by a "C" or "F" (or "c" or "f")
- **Problem 4: Obtain a temperature. If it is in celsius, return the number in fahrenheit; if in fahrenheit, return the number in celsius.**
 - Number can be an integer or a float, positive or negative
 - There must be only one number in the string
 - Character "C" or "F" implies what we are converting from and to.

Input validation plus more

```
#!/usr/bin/env python3

import re

input = input("Enter a temperature: ")
input.rstrip("\n")

matchobj = re.search(r"^-?([0-9]+(\.[0-9]*))?\s*([CF])$", input, re.IGNORECASE)
if matchobj:
    input_num, _, type = matchobj.groups()
    input_num          = float(input_num)

    if type == "C" or type == "c":
        celsius      = input_num
        fahrenheit   = (celsius * 9 / 5) + 32
    else:
        fahrenheit = input_num
        celsius     = (fahrenheit - 32) * 5 / 9

    print ("%.2f C is %.2f F\n" % (celsius, fahrenheit))
else:
    print ('Expecting a number followed by "C" or "F",')
    print ('so I cannot interpret the meaning of', input)
```

Notice how we indicate that case is to be ignored. The `re` module contains a large number of these kinds of options.

Problem Solving

- Our solution to Problem 4 still has some flaws
 - Cannot enter a number less than one without a leading zero.
 - No leading spaces are permitted (i.e., we have general whitespace issues)
 - We are using [0-9] instead of \d
 - etc. etc.
- There are many ways to "skin" a regular expression
 - The lesson so far, however, is that coming up with a full regular expression for these kinds of matches can be an iterative process.
 - Must also be aware of how a language deals with metasymbols within strings (e.g., Perl and Ruby are a bit different than Python)